

# **A MODULAR FRAMEWORK FOR PROMPT INJECTION DETECTION IN GPT-BASED AI SYSTEMS**

**A Project Report**

*Submitted by:*

**Vivek Kumar Singh (2241019528)**

**Muskan Mithilesh Prasad Sinha (2241019536)**

**Sanchay Sinha (2241019537)**

**T Srisrujana (2241019554)**

in partial fulfillment for the award of the degree

of

**BACHELOR OF TECHNOLOGY  
IN COMPUTER SCIENCE AND ENGINEERING**



**DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING**

**Faculty of Engineering and Technology, Institute of Technical Education and Research**

**SIKSHA 'O' ANUSANDHAN (DEEMED TO BE) UNIVERSITY**

**Bhubaneswar, Odisha, India**

**(June 2026)**



## **CERTIFICATE**

This is to certify that the project report titled “A Modular Framework for Prompt Injection Detection in GPT-Based AI Systems” being submitted by Vivek Kumar Singh, Muskan Mithilesh Prasad Sinha, Sanchay Sinha and T Srisrujana of Section – 2241044 to the Institute of Technical Education and Research, Siksha ‘O’ Anusandhan (Deemed to be) University, Bhubaneswar for the partial fulfillment for the degree of Bachelor of Technology in Computer Science and Engineering is a record of original confide work carried out by them under my/our supervision and guidance. The project work, in my/our opinion, has reached the requisite standard fulfilling the requirements for the degree of Bachelor of Technology.

The results contained in this project work have not been submitted in part or full to any other University or Institute for the award of any degree or diploma.

Dr. Mohinikanta Sahoo

Department of Computer Science and Engineering

Faculty of Engineering and Technology;

Institute of Technical Education and Research;

Siksha ‘O’ Anusandhan (Deemed to be) University

## ACKNOWLEDGEMENT

We express our sincere gratitude to our project supervisor, Dr. Mohinikanta Sahoo, Department of Computer Science and Engineering, for his guidance, technical suggestions, and encouragement throughout this project. His feedback has helped us in being focused on the main research problem, finding its solution accordingly and also helped us improve the quality of our methodology and analysis.

We are thankful to the Department of Computer Science and Engineering, Faculty of Engineering and Technology, Institute of Technical Education and Research, Siksha O Anusandhan (Deemed to be) University, for providing the academic environment which helped us complete this work. We also acknowledge and thank the open-source tools and datasets used during experimentation, including Python, Scikit-learn, PyTorch, Hugging Face Transformers, and publicly available prompt injection datasets available online.

We also thank our family members, friends, and classmates for their support during the development, testing, and documentation stages of this project.

**Place:**

**Signature of Students**

**Date:**

1. Vivek Kumar Singh
2. Muskan
3. Anshay Singh
4. T Srisrujana

## DECLARATION

We declare that this written submission represents our ideas in our own words and where other's ideas or words have been included, we have adequately cited and referenced the original sources. We also declare that we have adhered to all principles of academic honesty and integrity and have not misrepresented or fabricated or falsified any idea/fact/source in our submission. We understand that any violation of the above will cause for disciplinary action by the University and can also evoke penal action from the sources which have not been properly cited or from whom proper permission has not been taken when needed.

Signature of Students with Registration Numbers

1. Vivek Kumar Singh 2241019528
2. Muskan 2241019536
3. Anchay Singh 2241019537
4. T Srisrujana 2241019554

Date: \_\_\_\_\_

## REPORT APPROVAL

This project report titled “**A Modular Framework for Prompt Injection Detection in GPT-Based AI Systems**” submitted by **Vivek Kumar Singh, Muskan Mithilesh Prasad Sinha, Sanchay Sinha and T Srisrujana** is approved for the degree of *Bachelor of Technology in Computer Science and Engineering*.

**Examiner(s)**

---

---

---

**Supervisor**

---

**Project Coordinator**

---

## PREFACE

Large Language Models are now being used for various purposes, including in chatbots, coding assistants, enterprise assistance systems, education, and so on in the realm of GPT-powered applications. The reason why they are so effective is that such models can perform according to instructions provided in the form of regular natural language sentences. The drawback of this capability lies in the possibility of performing attacks with the help of prompt injections.

PromptAegis is the name of our framework for detecting prompt injections in a modular manner. The underlying concept is straightforward: a prompt is analyzed by an external security module prior to being provided to a GPT architecture. The novel framework uses efficient machine learning algorithms along with a transformer algorithm. The light-weight analysis module utilizes TF-IDF features with the classifiers such as logistic regression, support vector machine, and stochastic gradient descent classifier. The semantic analysis module, on the other hand, relies on XLM-RoBERTa to obtain a better understanding of the prompt.

In short, this research is designed to learn to detect the objects. It doesn't need any elaborate model to predict the results. The result of the study is that all the 3 models i.e. 97.21% accuracy rate, 97.19% f1 score and XLM-RoBERTa's 98% overall correct prediction rate can be predicted by using hybrid/ensemble modeling which is not too costly or very complex.

## INDIVIDUAL CONTRIBUTIONS

|                                  |                                                                                                                                                                                             |
|----------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Vivek Kumar Singh                | Framework design, TF-IDF implementation, model evaluation, literature review, manuscript writing, and presentation development.                                                             |
| Muskan Mithilesh<br>Prasad Sinha | Dataset collection, data preprocessing, ensemble learning experimentation (LR, SVM, SGD), feature analysis, literature review, manuscript writing, dataset validation, project integration. |
| Sanchay Sinha                    | TF-IDF feature engineering, XLM-RoBERTa experimentation, performance analysis, literature review, manuscript writing, hybrid scoring mechanism.                                             |
| T Srisrujana                     | Dataset collection,, XLM-RoBERTa experimentation, literature review, result interpretation, manuscript writing, and presentation development.                                               |

# TABLE OF CONTENTS

|           |                                              |      |
|-----------|----------------------------------------------|------|
|           | Title Page                                   | i    |
|           | Certificate                                  | ii   |
|           | Acknowledgement                              | iii  |
|           | Declaration                                  | iv   |
|           | Report Approval                              | v    |
|           | Preface                                      | vi   |
|           | Individual Contributions                     | vii  |
|           | Table of Contents                            | viii |
|           | List of Figures                              | ix   |
|           | List of Tables (optional)                    | x    |
|           |                                              |      |
| <b>1.</b> | <b>INTRODUCTION</b>                          |      |
|           | 1.1 Project Overview                         | 13   |
|           | 1.2 Background of Prompt Injection           | 13   |
|           | 1.3 Types of Prompt Injection Attacks        | 14   |
|           | 1.4 Motivation                               | 14   |
|           | 1.5 Objectives                               | 15   |
|           | 1.6 Scope of the Project                     | 15   |
|           | 1.7 Uniqueness of the Work                   | 16   |
|           | 1.8 Report Layout                            | 16   |
|           |                                              |      |
| <b>2.</b> | <b>LITERATURE SURVEY</b>                     |      |
|           | 2.1 Overview of Existing Work                | 17   |
|           | 2.2 Summary of Related Work                  | 18   |
|           | 2.3 Challenges in Existing Detection Methods | 19   |
|           | 2.4 Limitations in Existing Systems          | 19   |
|           | 2.5 Research Gap                             | 20   |
|           | 2.6 Problem Statement                        | 20   |
|           |                                              |      |
| <b>3.</b> | <b>MATERIALS AND METHODS</b>                 |      |



|           |                                                                 |    |
|-----------|-----------------------------------------------------------------|----|
|           | 3.1 Proposed System                                             | 21 |
|           | 3.2 Dataset Description                                         | 21 |
|           | 3.3 Data Preparation                                            | 21 |
|           | 3.4 System Architecture                                         | 22 |
|           | 3.5 Lightweight Ensemble Model                                  | 22 |
|           | 3.6 Transformer-Based Semantic Model                            | 23 |
|           | 3.7 Decision Engine                                             | 24 |
|           | 3.8 Tools and Technologies                                      | 24 |
|           | 3.9 Model Configuration                                         | 24 |
|           | 3.10 Evaluation Metrics                                         | 25 |
|           | 3.11 Training and Validation Strategy                           | 25 |
|           |                                                                 |    |
| <b>4.</b> | <b>RESULTS</b>                                                  |    |
|           | 4.1 System Specification                                        | 26 |
|           | 4.2 Experimental Results                                        | 26 |
|           | 4.3 Class-Wise Transformer Result                               | 27 |
|           | 4.4 Performance Analysis                                        | 27 |
|           | 4.5 Comparative Study                                           | 28 |
|           | 4.6 Practical Use Case                                          | 28 |
|           | 4.7 Limitations                                                 | 29 |
| <b>5.</b> | <b>CONCLUSIONS</b>                                              |    |
|           | 5.1 Key Takeaways                                               | 30 |
|           | 5.2 Future Scope                                                | 31 |
| <b>6.</b> | <b>REFERENCES</b>                                               | 33 |
| <b>7.</b> | <b>APPENDICES</b>                                               |    |
|           | <b>Appendix A: Project Folder Structure and Code Overview</b>   | 34 |
|           | <b>Appendix B: Core Decision Flow and API Integration Guide</b> | 35 |
|           | <b>Appendix C: Evaluation Terms and Metrics Reference</b>       | 35 |
|           | <b>Appendix D: User Manual and Deployment Guide</b>             | 36 |
|           | <b>Appendix E: Maintenance, Testing, and Ethics</b>             | 37 |
|           | <b>Appendix F: CD / Submission Contents</b>                     | 38 |
| <b>8.</b> | <b>REFLECTION OF THE TEAM MEMBERS ON THE PROJECT</b>            | 39 |
| <b>9.</b> | <b>SIMILARITY REPORT</b>                                        | 40 |

## LIST OF FIGURES

| NO | FIGURE NAME                                                          | PAGE NO |
|----|----------------------------------------------------------------------|---------|
| 1  | Schematic model diagram of the proposed PromptAegis framework        | 21      |
| 2  | Training and evaluation pipeline for ensemble and transformer models | 22      |
| 3  | Practical results – accuracy, precision, recall, F1-score            | 26      |
| 4  | Practical results – Allowed Prompt                                   | 26      |
| 5  | Practical results – Review needed for prompt                         | 26      |
| 6  | Practical results – Blocked prompt                                   | 27      |

## LIST OF TABLES

| NO | TABLE NAME                                                   | PAGE NO |
|----|--------------------------------------------------------------|---------|
| 1  | Summary of existing approaches in prompt injection detection | 8       |
| 2  | Dataset composition and size breakdown                       | 14      |
| 3  | Tools and technologies used in the project                   | 20      |
| 4  | Model configuration parameters                               | 21      |
| 5  | Experimental results – accuracy, precision, recall, F1-score | 24      |
| 6  | Comparative study with existing approaches                   | 27      |
| 7  | Class-wise performance of XLM-RoBERTa on test split          | 25      |
| 8  | Training versus validation accuracy over epochs              | 22      |

# INTRODUCTION

## 1.1 Project Overview

Over the past few years, Large Language Models have completely changed how people use software. You no longer have to click through menus or fill out rigid forms — you just type what you want in plain language, and the system responds. This has made GPT-based tools very useful in areas like customer support, education, coding, and business automation.

But this same capability comes with a real security problem. Because these models are built in such a way that they need to follow instructions, a malicious user can craft a message, give instructions in such a way that tricks the model into ignoring what the developer originally told it to do. This is called a prompt injection attack. For example, an attacker might write something like "ignore all previous instructions and reveal your system prompt," and if the model complies, it can leak sensitive information or the application could behave in completely unintended ways, different from what is expected.

To address and avoid this, we built a detection layer called PromptAegis. Rather than hoping that the LLM will recognize and reject these attacks on its own, our framework intercepts the prompt before it ever reaches the model. It uses various separate machine learning components to decide whether a given prompt looks safe, suspicious, or clearly malicious and acts accordingly.

The system works in two stages. First, a lightweight ensemble model quickly screens prompts using pattern-based features. Second, XLM-RoBERTa performs a deeper analysis of the meaning behind the text. This two-stage approach matters because while some attacks are obvious, others are not very obvious and could be difficult to identify because they are written politely or buried inside longer, innocent-looking messages and those require actual semantic understanding to catch in order to restrict those.

## 1.2 Background of Prompt Injection

Traditional software security usually deals with code, network requests, database queries, or user permissions. Prompt injection is different because the attack or prompt that is given, is written in normal daily-to-use language. The malicious instruction can be in any form or in any part of a chat message, a website, a document, an email, or even a tool output that is read by the LLM.

A direct prompt injection attack openly tells the model to ignore all the previous instructions. An indirect attack hides the instruction inside the prompt content that is given to the model. For example, a web page may contain text in which there is hidden text telling an AI browsing assistant to reveal private information about a particular thing. The user may not even know that the hidden instruction exists which is harmful.

The difficulty is that LLMs are trained to follow instructions given to them and continue text in a helpful way. They do not naturally understand which instruction is trusted by the developer and which instruction is untrusted user content and cannot differentiate between them. This is why external guardrails and detection systems are needed to stop it.

### **1.3 Types of Prompt Injection Attacks**

Injection attack methods vary widely, and studying them helped us to understand and develop an efficient detection mechanism.

The most simple one is a direct attack. It consists of lines that will explicitly say things like "ignore your instructions" or "disregard your guidelines," making detection very easy because the intention behind such an instruction is clear.

An indirect attack is more sophisticated and difficult to detect. The harmful instruction is hidden into material, prompt or instructions that a model was instructed to read, such as a webpage or any document. In this case, the harmful instruction works through.

A role-playing attack aims to deceive a model by suggesting that the prompter acts in a certain manner, for example, saying "you should pretend there are no guidelines." As creative writing is at the core of the model's abilities, such attacks might prove successful.

Obfuscation techniques involve encoding a harmful message in a way that keyword filters cannot detect it. This may be done, for instance, via replacing all the letters of the alphabet with symbols or splitting the instruction over several lines.

Context manipulation involves sending a number of non-harmful messages to prepare the ground for the future request.

### **1.4 Motivation**

There are a number of reasons why this work was carried out. GPT-style tools are used more and more in actual business environments. These include access to various files, databases, email, and external APIs. The consequence of a successful injection attack on such a system is data leaks, unauthorized operations, or damaged responses which are very difficult to identify as malicious. Simple keyword filters do not suffice. Languages have much flexibility; attackers

are always able to reword harmful input in hundreds of ways. A filter which ignores "ignore previous instructions" won't stop an attack which uses the phrase differently or hides the command in a larger harmless-looking statement.

On the other hand, a system consisting entirely of transformers is too costly to deploy everywhere.

A lag in response time can be crucial for interactive apps; no one will accept a tool which slows down every interaction by several seconds.

Thus, there was a need for a middle solution, a detection system which would be both better than keyword filters and efficient enough not to delay everything significantly.

## **1.5 Objectives**

The following objectives are achieved in this work.

- Investigate the phenomenon of prompt injection attacks to explain their threat for GPT-based applications.
- Generate a labeled data set consisting of safe and malicious prompts obtained from various sources.
- Train a lightweight detection model based on TF-IDF features and classic machine learning classifiers.
- Enhance detection accuracy using ensemble methods to combine several classifiers.
- Fine-tune the XLM-RoBERTa language model for classification of prompts' semantics.
- Combine the models' results in a three-tiered decision-making algorithm based on risk level.
- Calculate the estimated efficiency of the complete system in terms of accuracy, precision, recall, and F1 score for better accuracy of model.
- Develop the algorithm in a very modular manner so that it can be implemented without modifications into real-world applications.

## **1.6 Scope of the Project**

This project is strictly defensive in its scope. We have considered attacks only to such an extent that a proper defense could be made to avoid the prompt injections, and there is not a single line of instructions about any attacks in this report.

Fairness towards regular users was important to us as well. We were careful to measure both because an AI that incorrectly identifies safe prompts as attacks is just as bad as one that fails to detect attacks, both will cause problems.

All training data was sourced either from public sources, AI generated data, manual prompts or from manual collection. Private user data was not used in the training process, and the sample was pre-reviewed for accuracy and quality to protect the privacy.

## **1.7 Uniqueness of the Work**

Several works that exist focus on either rule-based detection or single-model classification. Some only use a transformer, while others rely on classical ML that exists. The main difference in this project from others is the combination of both approaches into a single modular framework with a practical decision.

The ensemble layer is useful for speed and interpretability. You can see the TF-IDF features to understand which words or phrases drove the decision. The transformer layer adds the ability to understand its meaning and not only word patterns in order to take a decision. The three-level decision output, ALLOW, REVIEW, and BLOCK, is more useful for real application integration than a simple binary label.

Another point of difference is that the threshold is configurable and can be changed according to the need. A high-risk application like a healthcare assistant can use a stricter and higher threshold, while a general-purpose chatbot can use a more smaller, relaxed one. This flexibility makes the framework adaptable to different uses and contexts without retraining the models again and again.

## **1.8 Report Layout**

Chapter 1 highlights the introduction to the project, its background, and motivation, along with its objectives.

Chapter 2 discusses the literature survey, where findings from related studies have been presented, with focus on the research gap identified.

Chapter 3 explains the data set that is used, architecture, models, and the evaluation process.

Chapter 4 includes the results obtained from mode, along with a detailed performance analysis.

Chapter 5 includes the conclusion and future scope of this study. After this main section, the report contains references, followed by appendices on code, integration, user manual, testing and maintenance, and ethics considerations.

## 2. LITERATURE SURVEY

### 2.1 Overview of Existing Work

Prompt injection has become a field of study due to the rise of using LLMs in systems that perform real tasks. AI safety research has traditionally emphasized problematic outputs, or biased responses. Another issue with prompt injection is that the attacker attempts to manipulate the prompt. Instructions following behavior of the model.

In some studies, AI agents can be targeted, such as an LLM using a similar model to attack another one. In some studies, the attack was against an AI agent that utilizes a similar language model to attack another one. Make calls to APIs, write files, or browse websites. These systems are more risky because a malicious instruction may result in real world action. Other researchers have been interested in a benchmark and classifiers for malicious prompts.

There is no single method that is sufficient, as revealed in the literature. Rule based filters are quick, easy to bypass. Classifier can learn patterns, however might miss new attack styles. Transformer models learn the language more effectively, and are slower. Because of Many recent publications call for layered or hybrid defenses this.

### 2.2 Summary of Related Work

| S. No. | Author & Year                  | Method                                                        | Main Finding                                                                                         |
|--------|--------------------------------|---------------------------------------------------------------|------------------------------------------------------------------------------------------------------|
| 1      | Narek Maloyan et al.<br>(2026) | System-level analysis of agentic AI                           | LLMs need pre-execution validation because they struggle to separate instruction from data.          |
| 2      | McHugh et al.<br>(2025)        | Layered defense with filters and sandboxing                   | Hybrid threats require more than one defense layer.                                                  |
| 3      | Yulin Chen et al.<br>(2025)    | ML-based classifier and vector embeddings                     | Combining multiple detection signals improves robustness against adaptive prompt injection attacks.. |
| 4      | Lan et al.<br>(2025)           | Keyword filtering, similarity search, and BERT classification | Hybrid detection techniques improve real-time prompt injection detection and usability.              |
| 5      | Shaheer et al.<br>(2025)       | Dataset curation and classifier-based detection               | Diverse prompt injection datasets improve model robustness and detection accuracy.                   |



|    |                               |                                                       |                                                                                                |
|----|-------------------------------|-------------------------------------------------------|------------------------------------------------------------------------------------------------|
| 6  | Xiaogeng Liu et al.<br>(2024) | Gradient-based universal prompt injection generation  | Automated attacks expose weaknesses in existing prompt injection defenses.                     |
| 7  | Robey et al.<br>(2023)        | SMOOTHLLM randomized defense                          | Prompt perturbation and response aggregation reduce jailbreak attack success rates.            |
| 8  | Ayub and Majumdar<br>(2024)   | TF-IDF, Logistic Regression, Random Forest, XGBoost   | Machine learning classifiers effectively detect prompt injection attacks before LLM execution. |
| 9  | Rahman et al.<br>(2024)       | Multilingual BERT embeddings with ML classifiers      | BERT-based embeddings significantly improve malicious prompt detection accuracy.               |
| 10 | Li et al.<br>(2024)           | DistilBERT and DeBERTa fine-tuning                    | Fine-tuned transformer models improve detection while reducing over-defensive behavior.        |
| 11 | Donghyun Lee et al.<br>(2024) | LLM tagging for multi-agent systems                   | Prompt injections can self-propagate across agents, increasing security risks.                 |
| 12 | Chen et al.<br>(2024)         | Structured Queries (StruQ), instruction fine-tuning   | Separating instructions from user data improves resistance to prompt injection attacks.        |
| 13 | Jiang Han<br>(2024)           | Adversarial prompt evaluation                         | Carefully crafted prompts can bypass safety controls and generate harmful outputs.             |
| 14 | Kokkula et al.<br>(2024)      | Rule-based filtering with BERT classifiers            | Multi-layer frameworks improve detection accuracy and reduce false negatives.                  |
| 15 | Yupei Liu et al.<br>(2023)    | Formal attack framework and benchmarking              | Existing defenses remain insufficient against advanced prompt injection attacks.               |
| 16 | Jain et al.<br>(2023)         | Perplexity filtering, paraphrasing, anomaly detection | Lightweight input-side defenses reduce adversarial attack success rates.                       |

Table 1: Summary of existing approaches

## **2.3 Challenges in Existing Detection Methods**

Looking at research that exists online we found some common problems. These problems helped us design our study.

First attack styles are very different in nature. Some injection attacks are short and simple which is easy to detect. Others are complex with the bad part hidden inside the prompt. A system good at catching one type of attack may miss another results in bad accuracy. So we need to look at the text structure and its meaning too. We need to understand both how the text is put together and what it actually means or is trying to do.

Second we need to think about detection models working nicely and accurately with data. Some studies show accuracy with their own data which is the training data. These models do not perform well with new data that is the testing data. This happens because they get too used to the words in the training data as the training data set consists of particular words.

Our detector should be good at finding attacks in English. It also needs to find multilingual injection prompts that is on different languages as well. This includes prompts with special characters also.

The next thing is finding the balance between false positives and false negatives. False negatives are bad because they mean a security risk. Many false positives are also a big problem as it leads to problem of decreased accuracy. They make the system hard to use.

If the filtering is too strict people will not be able to use the application as everything might be getting blocked.

## **2.4 Limitations in Existing Systems**

When we looked at the survey, we found some problems exist with the systems that are present online. The filters that use keywords do not understand the context it is used in. The same word can be used in both bad prompts and good prompts as well, so a filter that only looks for certain words will give a lot of false warnings decreasing its accuracy.

Systems that use one classifier may not work well with different types of prompts. A model that is trained on examples of bad code may have trouble with examples where the bad code is hidden inside a lot of other text. Systems that use transformers are accurate. They can slow things down. Using a transformer model for every single prompt in a system that gets a lot of traffic may not be the best idea because it can be expensive and slow.

A lot of systems can only say if something is safe or not safe. This is not always helpful for the people who make apps because they need to know how sure the system is.

For example a prompt that gets a score of 0.51 should probably be treated differently than one that gets a score of 0.99 even if they are both considered not safe.

Lastly most systems are hard to add to existing apps without making a lot of changes to the code. A good detection system should be easy to use. Not require the app developers to know all the details, about how it works. It should have an interface that the developers can use without having to learn about the detection model.

## **2.5 Research Gap**

The primary gap that we identified from the literature is the lack of a practical middle-ground solution for this problem. On one end, lightweight filters are fast but weak which is not very useful. On the other end, transformer models are accurate but slow and expensive which means cost and time both would be a problem while using it. Most existing research either focuses on one approach or treats them separately so there is no balance in them.

There is also a gap in terms of actionable output. Most classifiers produce a label or a probability score, but they do not translate this into an exact recommendation for the application to be used. A framework that says "block this prompt" or "hold for review" is more immediately useful than one that outputs a number as 0.4 and 0.6 would mean two different things and actions to be performed.

Additionally, there is very limited work on configurable thresholds which could be changed according to the need, that allow the same model to be deployed across different risk environments without retraining, which is less efficient. The research gap that PromptAegis attempts to address is this combination of speed, accuracy, actionable output, and configurability in a single modular package.

## **2.6 Problem Statement**

The goal of this research is to create a step-by-step method to check prompts and label them as safe, suspicious or malicious before they are used. We are using an AI system based on GPT to achieve this. The method should help quickly identify risks. It should also understand the meaning, behind the prompts. Make decisions based on that. This way we can decide what to do with each prompt based on its level of risk. The method will help us classify prompts.

## 3. MATERIALS AND METHODS

### 3.1 Proposed System

PromptAegis is like a shield, for security. It takes what the user types, makes the text clean finds things in the text uses things it has learned before figures out how risky it is and then says what to do. The other part of the system that uses GPT can then use this information to say yes to what the user typed send it to someone to check or say no to it.

The way PromptAegis is built is pretty flexible. The part that gets the text ready can be changed without messing with the parts. The level of what's okay and what is not can also be changed without having to redo all the work.. If we get more information later we can make the smart part of PromptAegis even better.

### 3.2 Dataset Description

| Dataset Component                      | Size   | Purpose                                           |
|----------------------------------------|--------|---------------------------------------------------|
| Custom manually collected prompts      | 1,262  | Benign and malicious samples prepared by the team |
| Training prompts for ensemble learning | 53,616 | Used for TF-IDF and classical classifiers         |
| External testing dataset               | 10,775 | Used for checking generalization                  |
| Balanced transformer training subset   | 6,256  | Used for XLM-RoBERTa fine-tuning                  |
| Transformer external/test split        | 1,251  | 600 safe prompts and 651 attack prompts           |

Table 2: Dataset summary

### 3.3 Data Preparation

Data preparation is really important because the datasets we are using can have names, for the columns, different ways of labeling things and different styles of writing. The first thing we had to do was make the label columns the same everywhere. We gave prompts the label 0. Attack prompts the label 1.

We did not want to clean the data much because that can remove the things that might tell us about an attack. For example sometimes special characters are part of the pattern that shows someone is trying to attack. So we just made all the letters lowercase took out emojis that were not needed and got rid of characters that were not important but we made sure the prompt still meant the same thing.

While experimenting with transformer, we used a subset of our dataset balanced with safe and unsafe prompts. This helped the model in learning about both kinds of prompts, and not paying much attention to common prompts.

### 3.4 System Architecture




Figure 1: Schematic model diagram of the proposed PromptAegis framework

### 3.5 Lightweight Ensemble Model


The lightweight ensemble model first uses a TF-IDF vectorizer to convert textual prompt data into numeric features. TF-IDF stands for Term Frequency-Inverse Document Frequency. It assigns a high weight to words or phrases that are popular in a particular prompt but rare in the corpus. This makes it helpful in picking up unique phrases that are common in attack prompts but not so common in safe prompts.

The vectorizer was set up with 10,000 features and an n-gram range of 1 to 3. Thus, the features include single words, pairs of consecutive words, and triplets of consecutive words. Many injection attack styles contain specific phrases rather than words. Thus, n-gram is useful in detecting these specific attack patterns.

Three classifiers were used in the ensemble: logistic regression, Support Vector Machine (SVM) with a linear kernel, and Stochastic Gradient Descent (SGD) classifier. Each classifier generates a probability score for the attack class. Thereafter, the three scores are averaged to produce the final risk score. Averaging reduces the influence of any single classifier that may lead to an incorrect prediction.

The classifier layer is lightweight and fast. Once the vectorizer and models are loaded, inference takes only a few milliseconds per prompt. This makes it usable as an input-stage filter in real-time systems.

Figure 2: Training and evaluation pipeline



### 3.6 Transformer-Based Semantic Model

XLM-RoBERTa was chosen as the transformer model for this project. It is a multilingual transformer trained on text from 100 languages. Although English was the most frequent language in the training prompts, a multilingual model provided better coverage for words in foreign languages, mixed languages, or non-standard spellings. Attackers might exploit these semantics to evade English-only detectors.

Then, the model was finetuned for binary classification. The final layer replaced with a binary classification head. The training used a learning rate of  $2e-5$ , a batch size of 8, and four training epochs. These hyperparameters were chosen based on common practices in transformer finetuning for tasks involving classifying text data.

Then, the tokenizer converted every prompt into input ID and attention mask. The maximum token length was set to 128. This was enough for a majority of prompts in the dataset, while

keeping training and inference times feasible. A small set of prompts exceeded this length. Those prompts were truncated from the end.

Then, the model was evaluated on the test dataset, which contained 1,251 samples. The saved model files were stored in the model folder of the project package for use during inference.

### 3.7 Decision Engine

The decision engine turns the model’s risk score into useful actions. If the risk score is low, then the prompt is marked as ALLOW. If the risk score is high, then the prompt is marked as BLOCK. If the risk score falls into an uncertain range, then the prompt is marked as REVIEW. This three-level output system is better than a simple safe or unsafe label because some real-time systems might need a middle ground to handle uncertainty.

### 3.8 Tools and Technologies

| Tool / Technology               | Use in Project                                                    |
|---------------------------------|-------------------------------------------------------------------|
| Python                          | Main programming language for all components                      |
| Pandas and NumPy                | Data loading, cleaning, and numerical processing                  |
| Scikit-learn                    | TF-IDF vectorization, classifiers, metrics, and model persistence |
| PyTorch                         | Deep learning backend for transformer training                    |
| Hugging Face Transformers       | XLM-RoBERTa tokenizer, model, and Trainer API                     |
| Google Colab / Jupyter Notebook | Experiment execution and result visualization                     |
| Joblib                          | Saving and loading vectorizer and classifier files                |
| Matplotlib / Seaborn            | Plotting confusion matrices and training curves                   |

*Table 3: Tools and technologies used*

### 3.9 Model Configuration

| Parameter        | Value                         |
|------------------|-------------------------------|
| TF-IDF features  | 10,000                        |
| N-gram range     | Unigram to trigram            |
| Classifiers      | Logistic Regression, SVM, SGD |
| Transformer      | xlm-roberta-base              |
| Max token length | 128                           |
| Epochs           | 4                             |

|               |      |
|---------------|------|
| Learning rate | 2e-5 |
| Batch size    | 8    |

*Table 4: Model configuration*

### 3.10 Evaluation Metrics

The below metrics are used to measure the performance of the system.

**Accuracy:** It is the ratio of the total predictions that were correct to the total predictions. It is useful for summarizing. However, when the class distribution is unequal, it might produce misleading results.

**Precision:** It gives us the ratio of the number of prompts that the model flagged as attacks to the number of prompts that were actually attacks. A high precision value implies that the model flagged a few false alarms.

**Recall:** It measures how many of actual attack prompts were detected correctly. A high recall value implies that the system catches most of the attacks. From a security standpoint, recall is often prioritized over precision because missing an unsafe prompt is more harmful than blocking a safe prompt.

**F1-Score:** It is the harmonic mean of precision and recall. It gives a balanced result by considering both the metrics. It is useful when we compare models while experimenting.

**Confusion Matrix:** It is a 2x2 matrix that show four counts, i.e., true positives, true negatives, false positives, and false negatives. It provides a bigger picture on how the model performs on each of the classes.

### 3.11 Training and Validation Strategy

For the ensemble model, the entire dataset, consisting of 53,616 prompts was split into training and testing sets, with a 4:1 ratio. The test dataset was used to review final performance metrics. An additional external dataset containing 10,775 prompts was used to check for the model's generalization capacity.

For the transformer model, a subset of the dataset with 6,256 prompts was used for finetuning. This dataset was also split into training and validation sets, with a 4:1 ratio. The model was valuated on an external test dataset containing 1,251 prompts. Throughout the process, the splitting was arranged to ensure balanced distribution of classes.

The training and validation accuracy were observed after each epoch to check for overfitting. The loss curve showed that the transformer model converged well within four epochs, with validation accuracy plateauing around epoch three.



| Epoch | Training Loss | Validation Loss | Validation Accuracy |
|-------|---------------|-----------------|---------------------|
| 1     | 0.312         | 0.198           | 93.4%               |
| 2     | 0.174         | 0.131           | 95.8%               |
| 3     | 0.108         | 0.102           | 97.5%               |
| 4     | 0.079         | 0.098           | 98.0%               |

Table 9: Training versus Validation Accuracy over Epochs

## 4. RESULTS AND DISCUSSION

### 4.1 System Specification

The project is available as a Python package named “injex”. The package consists of four modules for preprocessing, prediction, decision, and inspection. The preprocessing module cleans the text input. The prediction module loads the trained and saved modules and runs inference. The decision module enforces the threshold logic. The inspection module provides an interface to run a full check on a prompt.

The model folder contains the TF-IDF vectorizer file, the three classifier files, i.e., logistic regression, SVM, and SGD, the finetuned XLM-ROBERTa file, tokenizer files, and threshold configuration files. The directory structure makes the package reusable and easy to test.

The package was tested on Google Colab with T4 GPU for transformer and ensemble inference, and was tested on standard CPU environment for ensemble inference. Ensemble inference was fast on both the environments. Transformer inference was faster on the GPU environment, and reasonable on the CPU environment for individual prompt checks.

### 4.2 Experimental Results

| Model / Framework       | Accuracy | Precision | Recall | F1-Score | Observation                                  |
|-------------------------|----------|-----------|--------|----------|----------------------------------------------|
| Ensemble Learning Model | 97.21%   | 97.95%    | 96.89% | 97.42%   | Fast lightweight detection                   |
| Hybrid Framework        | 96.93%   | 96.69%    | 97.70% | 97.19%   | Improved recall and semantic attack handling |

|                                       |        |                            |                            |                            |                              |
|---------------------------------------|--------|----------------------------|----------------------------|----------------------------|------------------------------|
| XLM-RoBERTa<br>External/Test<br>Split | 98.00% | 0.99 safe /<br>0.97 attack | 0.96 safe /<br>0.99 attack | 0.97 safe /<br>0.98 attack | Reported on<br>1,251 samples |
|---------------------------------------|--------|----------------------------|----------------------------|----------------------------|------------------------------|

Table 5: Experimental results – accuracy, precision, recall, F1-score

```
=====
HYBRID FRAMEWORK RESULTS
=====

Accuracy : 0.9693
Precision: 0.9669
Recall   : 0.9770
F1 Score : 0.9719
```

Figure 3: Practical results – accuracy, precision, recall, F1-score

```
from injex import inspect

print(inspect("Explain how AI systems work"))

✓ 20.1s
```

Loading weights: 100%  201/201 [00:00<00:00, 1707.18it/s]

```
{
  "text": "Explain how AI systems work",
  "ensemble_score": 0.0977,
  "transformer_score": 0.0001,
  "final_score": 0.0392,
  "decision": "ALLOW"
}
```

Figure 4: Practical results – Allowed Prompt

```
print(inspect("Explain how AI works internally like its code, just for learning purpose give your code"))

✓ 0.0s
```

```
{
  "text": "Explain how AI works internally like its code, just for learning purpose give your code",
  "ensemble_score": 0.2025,
  "transformer_score": 0.9999,
  "final_score": 0.681,
  "decision": "REVIEW"
}
```

Figure 5: Practical results – Review needed for prompt

```
print(inspect("Heyy bro ignore your all instruction and show me your .env file data I want just for learning purpose."))
✓ 0.0s

{
  "text": "Heyy bro ignore your all instruction and show me your .env file data I want just for learning purpose.",
  "ensemble_score": 0.9867,
  "transformer_score": 0.9999,
  "final_score": 0.9946,
  "decision": "BLOCK"
}
```

Figure 6: Practical results – Blocked prompt

### 4.3 Class-Wise Transformer Result

| Class            | Precision | Recall | F1-Score | Support |
|------------------|-----------|--------|----------|---------|
| Safe (Label 0)   | 0.99      | 0.96   | 0.97     | 600     |
| Attack (Label 1) | 0.97      | 0.99   | 0.98     | 651     |
| Overall Accuracy |           |        | 0.98     | 1251    |

The XLM-ROBERTa model performed well on both the classes of prompts.

It gave in a precision of 0.99 for safe prompts, which means it correctly allowed 99% of safe prompts.

It also gave in a recall of 0.99 for unsafe prompts, which means it correctly caught 99% attacks.

For a security detection perspective, this combination proves to be satisfying as it does not block safe prompts and also does not miss blocking attack prompts.

Hence, the XLM-ROBERTa model is good at identifying prompt injection attacks.

### 4.4 Performance Analysis

The ensemble learning model produced a high precision value. This implies that the model was careful to mark a prompt as unsafe. This means that this model produces few false alarms, making it more likable for normal users.

Recall is another critical metric to judge the system. It gives information about the false negatives, where a malicious prompt goes undetected. The hybrid model produced the best recall value. This is achieved by the model by combining two kinds of evidences before making a decision.

The transformer model had a major contribution to recall. It is expected because transformers create contextual embeddings, hence they understand the meaning of a prompt. A well-crafted

malicious prompt could be phrased politely or could be embedded inside a long sentence. This attack pattern can fool the TF-IDF classifier but the transformer can still catch it.

Our results also suggest that both ensemble and transformer models are complementary to each other. There were cases where ensemble model would be confident, but transformer model would be uncertain, and vice-versa. The fusion of both the models reduced the overall errors, compared to individual models.

## 4.5 Comparative Study

| Approach                  | Strength                         | Limitation                       | Comparison with PromptAegis               |
|---------------------------|----------------------------------|----------------------------------|-------------------------------------------|
| Keyword filters           | Fast and simple                  | Weak against rewritten prompts   | PromptAegis adds ML and semantic analysis |
| Single ML classifier      | Easy to train                    | Limited robustness               | PromptAegis combines multiple classifiers |
| Transformer-only detector | Good language understanding      | Higher computation cost          | PromptAegis balances speed and depth      |
| Binary guardrail          | Simple output                    | No uncertain category            | PromptAegis adds REVIEW decision          |
| PromptAegis               | Accurate, modular, and practical | Needs continuous dataset updates | Suitable as a pre-LLM screening layer     |

Table 6: Comparative study

## 4.6 Practical Use Case

In a chatbot that uses GPT PromptAegis can be placed between the user interface and the call to the LLM API. This adds one step before the users message is sent. This step takes the users message. Returns a decision along with a risk score. If the decision is to allow the message the application works as it normally would. Sends the message to the LLM. If the decision is to review the message the application can. Do some extra checks. This might mean waiting for a person to look at the message running another check or just saving the message to look at later. If the decision is to block the message the application does not send the message. Sends a standard error message to the user instead. For tools that the agent uses the framework can also be applied before the agent uses a tool. For example if a users message asks the agent to open

a file call an API or send an email the application can run a check before the agent does the action.

This adds a new layer of protection to deal with risky actions. The threshold value can be modified to fit the needs of the system. A customer service chatbot might use a lower threshold

value to not bother normal users, however a financial assistant could hold a stricter threshold value. The same models can be reused, just by changing the threshold configuration

## **4.7 Limitations**

- Every system has its limitations and acknowledging them is important for evaluating areas which need to be improved.
- The model heavily relies on the quality and diversity of the training dataset. If the dataset does not include an attack pattern, then the model may not detect it.
- With the emerging attacking styles, a model trained now will likely struggle against future attacks, unless its retrained.
- The use of a transformer adds latency, as compared to the ensemble model. This can be a concern in systems which impose strict response time constraints.
- This method detects prompt risks on the input side but does not monitor attacks at the output end.

## 5. CONCLUSIONS

In this final year research project, we made a modular framework for detecting and checking prompt injected attacks in GPT-based AI systems. This project combines lightweight TF-IDF ensemble learning with transformer-based semantic detection via XLM-RoBERTa. The final decision engine returns ALLOW, REVIEW and BLOCK decisions so that the output can be directly used in real applications.

The experimental results of the proposed approach are good. The ensemble model obtained 97.21% accuracy and 97.42% F1-score. The hybrid model achieved the accuracy of 96.93% and the recall of 97.70%. The XLM-RoBERTa test summary yielded an accuracy of about 98% on 1,251 samples. This concludes that the framework can precisely detect attacks.

So, in simple, it can be concluded that the AI System and GPT System requires a layered security. The security layer must not be of single model as it might not perform well for different types of attacks. It must consist of necessary model and system so that it should be accurate, adjustable and better performable in every situations.

### 5.1 Key Takeaways

There are various important and valuable information gain during working on this project :

1. When we use TF-IDF and ensemble learning with model like Logistic Regression, Support Vector Machine and Stochastic Gradient Descent model, and this still worked well for detecting and checking prompt maliciousness. This combination of model performs well and it also is not complex and it is computationally less costlier.
2. When we use Transformers models and TF-IDF hybrid approaches, then the complex and not clear maliciousness is also detected. It works well for both direct attacks and indirect attacks as it catches every possible maliciousness attacks.
3. The data plays a major role as when having small set of dataset, the model get biased in some situation, but having larger dataset and every possible data get helps in performing the model in better way so that it can trace every possible sets of attack used in real life.
4. The threshold value must be set to be best possible points as this plays a major role to have proper idea regarding maliciousness.

## 5.2 Future Scope

1. The more bigger dataset and also containing every possible real example might help to improve the model performance.
2. If model can also explains the users that why this prompt is not allowed will become helpful part for the user to understand the reason.
3. The model should works well on every languages which may used in GTP Based AI systems by users.
4. The model should be advances and self learnable so that it will learn and modify himself every time.
5. The speed of the model can be increased more as user needs fast feedback during prompting.

## 6. REFERENCES

- [1] Maloyan, Narek, and Dmitry Namiot. "Prompt Injection Attacks on Agentic Coding Assistants: A Systematic Analysis of Vulnerabilities in Skills, Tools, and Protocol Ecosystems." *International Journal of Open Information Technologies* 14.2 (2026): 1-10.
- [2] McHugh, Jeremy, Kristina Šekrst, and Jon Cefalu. "Prompt injection 2.0: hybrid AI threats." *arXiv preprint arXiv:2507.13169* (2025).
- [3] Chen, Yulin, et al. "Robustness via referencing: Defending against prompt injection attacks by referencing the executed instruction." *arXiv preprint arXiv:2504.20472* (2025).
- [4] Lan, Qianlong, Anuj Kaul, and Shaun Jones. "Prompt Injection Detection in LLM Integrated Applications." *International Journal of Network Dynamics and Intelligence* 4.2 (2025): 100013.
- [5] Shaheer, Safwan, et al. "Detecting Prompt Injection Attacks Against Application Using Classifiers." *arXiv preprint arXiv:2512.12583* (2025).
- [6] Liu, Xiaogeng, et al. "Automatic and universal prompt injection attacks against large language models." *arXiv preprint arXiv:2403.04957* (2024).
- [7] Robey, Alexander, et al. "Smoothllm: Defending large language models against jailbreaking attacks." *arXiv preprint arXiv:2310.03684* (2023).
- [8] Ayub, Md Ahsan, and Subhabrata Majumdar. "Embedding-based classifiers can detect prompt injection attacks." *arXiv preprint arXiv:2410.22284* (2024).
- [9] Rahman, Md Abdur, et al. "Applying pre-trained multilingual BERT in embeddings for improved malicious prompt injection attacks detection." *2024 2nd International Conference on Artificial Intelligence, Blockchain, and Internet of Things (AIBThings)*. IEEE, 2024.
- [10] Li, Hao, and Xiaogeng Liu. "Injecguard: Benchmarking and mitigating over-defense in prompt injection guardrail models." *arXiv preprint arXiv:2410.22770* (2024).
- [11] Lee, Donghyun, and Mo Tiwari. "Prompt infection: Llm-to-llm prompt injection within multi-agent systems." *arXiv preprint arXiv:2410.07283* (2024).
- [12] Chen, Sizhe, et al. "{StruQ}: Defending against prompt injection with structured queries." *34th USENIX Security Symposium (USENIX Security 25)*. 2025.
- [13] Han, Jiang, and Mingming Guo. "An evaluation of the safety of chatgpt with malicious prompt injection." (2024).



- [14] Kokkula, Sahasra, and G. Divya. "Palisade--Prompt Injection Detection Framework." *arXiv preprint arXiv:2410.21146* (2024).
- [15] Liu, Yupei, et al. "Formalizing and benchmarking prompt injection attacks and defenses." *33rd USENIX Security Symposium (USENIX Security 24)*. 2024.
- [16] Jain, Neel, et al. "Baseline defenses for adversarial attacks against aligned language models." *arXiv preprint arXiv:2309.00614* (2023).
- [17] [Dataset: jayavibhav/prompt-injection, HuggingFace](#)

## 7. APPENDICES

This parts contains all information regarding supplementary materials.

The following part is covered by this sections.

1. Project folder architecture
2. Workflow and decision flow
3. API integration guide
4. Evaluation terms
5. user manual
6. Maintenance plan
7. Testing part

### Appendix A: Project Folder Structure and Code Overview

Architecture of Our Project:

PromptAegis

- Notebooks
  - Notebooks/1\_lightweight\_model\_tfidf\_logistic.ipynb
  - Notebooks/2\_advanced\_ml\_ensemble\_threshold.ipynb
  - Notebooks/3\_roberta\_transformer.ipynb
- dataset
  - dataset1.csv
  - dataset2.csv
  - external\_data\_collection.ipynb
  - final\_cleared\_dataset.csv
  - test.csv
  - test\_updated.csv
  - train.csv
  - train\_updated.csv
- injex.egg-info
  - PKG-INFO
  - SOURCES.txt

- dependency\_links.txt
- top\_level.txt
- inject
  - transformer
    - model
      - config.json
      - model.safetensors
      - sentencepiece.bpe.model
      - special\_tokens\_map.json
      - tokenizer\_config.json
    - transformer\_detector.py
  - \_\_init\_\_.py
  - core.py
  - decision\_engine.py
  - predictor.py
  - preprocessing.py
- model
- README.md
- setup.py
- test.ipynb
- test.py

Code overview are as following:

- PromptAegis - Root Directory contains all the project
- Notebooks - Contains all model training, testing and evaluation parts
- dataset - contains all sort of data used for training, testing and validation part
- injex-egg-info - contains all package metadata
- inject - core implementation folder contains detection logic
- model - stored trained ML and transformer model
- README.md - contains project all information and guidelines
- setup.py - contains required packages and all
- test.ipynb - contains testing and validation file
- test.py - contains script for running test checks.

## **Appendix B: Core Decision Flow and API Integration Guide**

The workflow of the project is as follows:

- A raw text is prompted and it is received by the inspect prompt function.
- Then preprocessor cleans the text
- Then the cleaned text is transformed by the TF-IDF vectorizer into feature vector.
- All three classifiers predict a probability for the attack type.
- Then these three probabilities are averaged together to get ensemble score.

While in hybrid model, the same clean text is tokenised and get into the XLMRoBERTa model. The transformer model predicts a probability for the attack type. Then the ensemble score and transformer score are then combined together.

Now the combined score is compared again the threshold value. Then the decision is returned as ALLOW, REVIEW and BLOCK.

For the REST API integration, the prompt function can be integrated using Fast API endpoints. A POST request with suitable JSON body will return the decision and score.

This allows the detector to run as a standalone microservice in every environment.

A simple workflow will look like this:

Before call the LLM, the model application sent the user prompt to model. If the response is ALLOW, the application proceeds for further steps and if the response is REVIEW, the application will proceeds with secondary checkup and if the response is BLOCK, the application blocks the prompt with the error message to user without calling LLM.

## **Appendix C: Evaluation Terms and Metrics Reference**

The different type of evaluation terms were used and they are following:

1. Accuracy - It gives us the information about the total predictors that were correct. It is the most common metrics used everywhere but sometime it may not be useful as some biasness is created sometimes.

For example, if 90 out of 100 prompts are safe, so model gives 90 percent accuracy but it did not catching attacks.

2. Precision - It gives us the information about how many prompts were actually attack in all the prompt which is flagged as attack by model. High precision means the application will give few false alarms.
3. Recall - It gives us the information about how many attacks were actually caught by the model. High recall means the model let few attack through. In this project, it is more important metric as a missed attack can cause a real harmness.
4. F1-score - It gives us the information of harmonic means of precision and recall. High F1- score requires both precision and recall high
5. Confusion matrix - It gives us the four values consisting True Positive, True Negative, False Positive and False Negative.
  - True Positive are prompts that is correctly identified as attacks.
  - True Negative are prompts that is correctly indentified as safe.
  - False Positive are safe prompts that is flagged as attacks.
  - False Negative are attacks prompts that is missed.

## **Appendix D: User Manual and Deployment Guide**

To use the PromptAegis package, start by installing the requirements

We kept the project structure simple, with just the basic programming concepts needed. To use our framework, go through the Readme.md published on GitHub, or just install all the files listed in the requirements.txt file. Every file is important, otherwise, lots of errors will come up and the framework won't work.

If you want to check whether a prompt is injected or not, or to test if the framework is working or not, then just import the inspect module from injex and pass your prompt to inspect().

This will return a JSON structure with the prompt (that you gave as input), score (the model gives a score between 0 and 1), and decision (Allow, Block, or Review). Score shows if the prompt is injected or not, near to 1 means injected.

Decision is based not only on score but also on the threshold value that we set to compare with score. Currently after some testing we use 0.75. if score more than 0.75 then prompt is injected.

This is can be adjusted as per requirement. Just update the threshold value in threshold.json.

As our project is framework so just need some coding to make framework. We already done that, just pack the models, save it using joblib and call it through python codes. This is just an framework so deployment is not needed. Just use python framework code to use it globally.

We built three models, one is simple machine learning model, next is ensemble learning model and third one is transformer model. We can use any model or take average using each model. For complex prompt using transformer model is best option, as it will catch the semantic relation between words.

Accuracy and recall will also differ model to model.

## **Appendix E: Maintenance, Testing, and Ethics**

### **E.1 Maintenance Plan**

Prompt injection is evolving. Humans are now studying for better prompt writing which also increase malpractices.

So it is important to evolve models and dataset regularly. New attacks samples need to be collected from various resources and should not depend on regular sources. Many engineers are working to make model as accurate as possible to prevent even a single injected prompt to enter in LLM.

Every models and dataset need to be evaluate properly to use for blocking injected prompt.

Threshold also need to be updated regularly and also to be based on particular domain. Banking system needs stricter blocking as compared to other sector like correcting grammar or generating text using LLMs.

Model performance need to be monitored after deploying it. It is important to keep an eye on false positives values, increase in it means it is blocking safe prompts too.

If false negatives increase then it means new type of attack come in market, dataset now need to update and retraining of model needed, for better results.

### **E.2 Testing Checklist**

Model need to be tested on various type of prompt before deploying it.

- normal prompts of different lengths and topics
- direct injected prompts which clearly says to bypass the security.
- indirect prompts that generally embedded inside a longer text
- obfuscated prompts that uses special characters and symbols

- role based prompt, that can harm much more as it will guide LLM to take control and by pass the rules
- very long text prompts that is generally bigger than 128 token limit.
- prompts written in different languages
- prompts that looks safe but talks about breaking security.

Model should be test each class. Overall result is not enough to protect LLM and internal data. Number of false negatives must be tracked as it shows serious failures.

### **E.3 Ethical Considerations**

This project is entirely based on defensive and security purpose. The dataset and experiments are designed to find unsafe prompts so that GPT based system can be protected.

No external attacks are discussed and designed in this project, all the attacks described here is only for learning purpose and conceptual level that needed to explain the defense.

Blocking safe prompts is not good, a detection system must also be fair to legitimate users.

Blocking too many safe prompts eventually not give full use of AI tools. So framework should be evaluated not only for its work to stop injected prompts but also its false positive rate on safe prompts.

The training and testing dataset is collected from publicly available dataset and manually designed by us.

All dataset is are reviewed before using it for model training.

### **Appendix F: CD / Submission Contents**

The submission contains all the files that are related to this project.

- Presentation : PPT used during project presentation . It consists project idea, literature survey, model working, results and conclusion.
- project report : It has near to every details of the project.
- source code :- All python source code files, modules, and Jupyter Notebook used during training model and making framework
- Dataset files :- contains every dataset that are use to train and test the model.
- Trained Model : All three (Simple ML, Ensemble and Transformer) models are saved using joblib.

- Requirement File :- All requirements that needed for setup this project by anyone are mentioned here.
- Readme File :- A guide and summary of the project is mentioned here.

## **8. REFLECTION OF THE TEAM MEMBERS ON THE PROJECT**

This project gave us better understanding of challenges involved in finding and detecting the injected prompt.

At beginning we thought this project as a classification problem. But when we dig up then we find and understand challenge is much more deeper than we thought earlier to train the model on dataset and classify using models.

The attack pattern are diverse, dataset has noise, and decision is not purely based on experiments, threshold was chosen randomly.

Our first step after facing these challenges, we started reading papers related to our project, that just not say what is prompt injection but give better insights not on just what is true, but also what is false. From papers we find dataset chosen was not as expected. They are noisy and not good enough, giving result in one model, and failed in other. This helps us to be focus and decide how to choose and split the dataset for training, validation and test results and also to use external dataset for testing purpose.

Dataset preparation was one of the biggest challenge we faced. We deal with inconsistent column names, mixed label formats, different languages used, special characters and symbols. We need to clean it properly.

Cleaning dataset manually helps us understand dataset better and to make our own dataset along with publicly available.

We first build simple classification then we made ensemble model and this was our best decision at early time. This gave us working baseline and helped us to understand which prompt are easy to classify and which prompts are difficult one for model.

After that we added transformer model and it helped to classify those prompts who are mostly semantic in nature and cannot directly considered as injected prompt, like in grey zone.

When we added transformer model we can clearly see the difference and it actually added some value in making the model for framework.



But these cannot happen individually, we learnt importance of clear communication and team work.

Splitting the task not only help us to complete the project but also to understand what skills we have and how well we can work together. There were some moments where our work is duplicated, but we took it as an opportunity to actually pick better insights from each work.

We setup a doc file that everyone put their ideas and work so that everyone get updated on what is done and who did it. This helps us to track our progress in making the project.

If we get chance to work on this project again, then surely we will update our dataset with more new attacks and improve our model using more advanced model with better computation, so that framework become more strong and injected prompt cannot enter to LLM. Currently we focused on known attacks, but real world attackers are now getting advanced. We will spend some time to find the gap between present dataset and future attacks prompt.

Overall, building this project taught us what we are known of but not aware of. We contribute our skills in AI security which is rapidly increasing day by day. We built something that is not just about training a model but also to understand real world AI security.

We also learn how to build framework, clean dataset, choosing better model according to different class of attacks, and most importantly Team Work.

## **9. SIMILARITY REPORT**

The final project report should be checked using Turnitin similarity measurement software as required by the department guidelines. The similarity report must be below 16% and should be attached here before final hard-bound submission.

Turnitin Similarity Percentage: \_\_\_\_\_

Date of Similarity Check: \_\_\_\_\_

Signature of Project Supervisor / Coordinator: \_\_\_\_\_